

Cloud Computing for HPC Simulations

Diego Montes <kabute@uvigo.es>

Goals

- Insights of HPC over Cloud Computing.
 - Give an starting point to start your own experiments in the Cloud.
-


What is the Cloud Anyway?

As defined by the NIST:

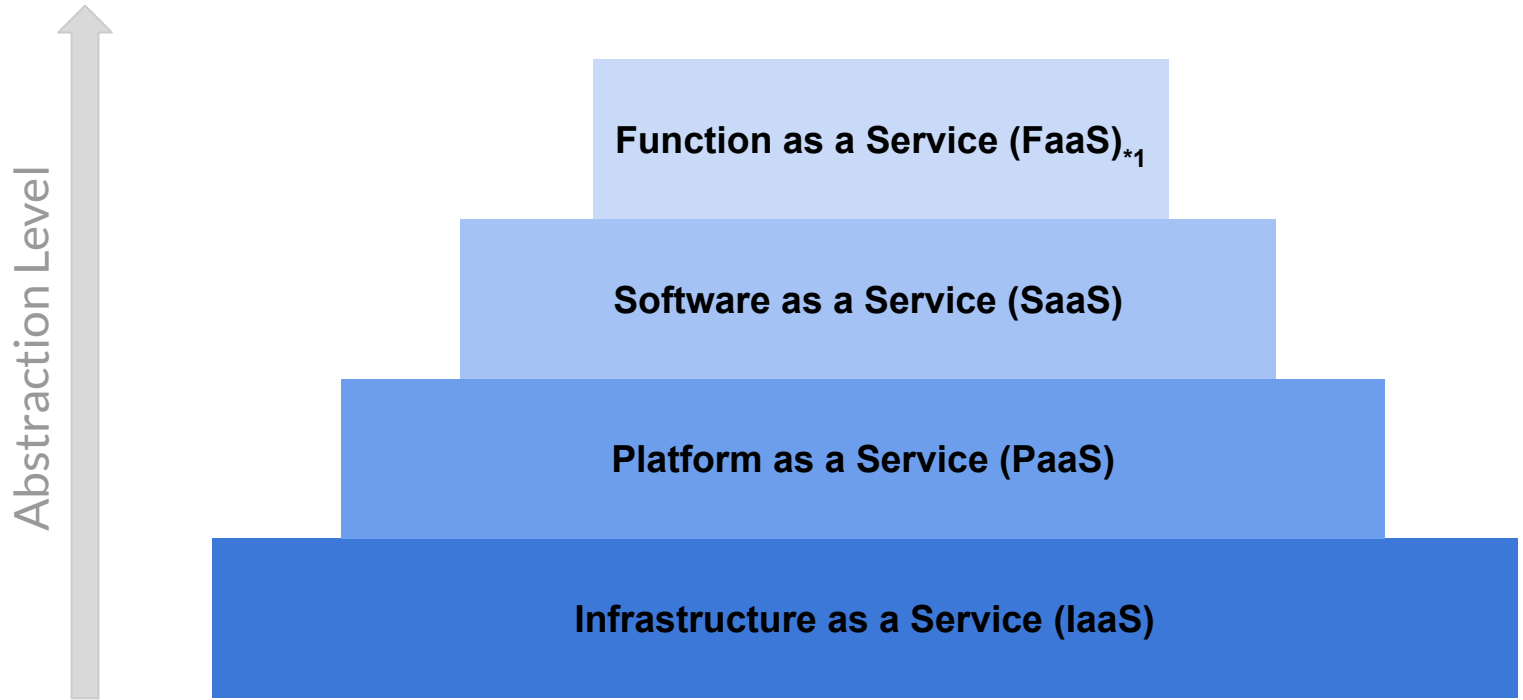
- On-demand Self-service
 - Broad Network Access
 - Rapid Elasticity
 - Measured Service
-

What is the Cloud Anyway?

A broad concept that encompasses the abstraction of computer systems (including their resources) and presents them in a ubiquitous way.



Cloud - Service Levels



*1: Including the concept of "Serverless"

Cloud - Service Levels

Rule of Thumb, choose wisely:

**Higher Abstraction ⇒ Less Operational Load ⇒
More Constraints ⇒ More Out-of-the-Box Features
⇒ Higher Vendor Dependency**

Is the Cloud Right for my Experiment(s)?

Maybe... Start by answering yourself:

If Legacy Code:

- Does to code run on commodity hardware?
 - How fault tolerant is the code?
-

Is the Cloud Right for my Experiment(s)?

If New Development:

- Can it take advantage (do I understand?) of the Cloud Paradigm? (APIs, ephemeral environments...)
-

Is the Cloud Right for my Experiment(s)?

In any case:

- What is the budget.
 - How many resources do you really need.
 - What are acceptable thresholds (e.g. failure/errors)
-

Developing HPC for the Cloud

Think **Cloud Native**:

A cloud native app is architected specifically to run in the elastic and distributed nature required by modern cloud computing platforms.

Developing HPC for the Cloud

- Immutable infrastructure, idempotence and containerization are not just a trend.

Developing HPC for the Cloud

- APIs are first class citizens.
 - Commodity hardware (reliable over unreliable).
 - Be prepared for failure: checking points, incremental backoff, different AZs...
-



Running HPC in the Cloud

Compared with a more traditional environment, **by default**, the Cloud is not very HPC “friendly”:

- No low latency networks.
 - Prone to different (points of) failure.
 - No limits***
-

Running HPC in the Cloud

- Scheduling is now transparent, (initially) no need for queues.
 - You might need to give a heads up to be sure resources are allocated.
-

Is the Cloud Right for my Experiment(s)?

But...

- Yes, there is HPC-friendly in the cloud (low latency, affinity, location...) ...
 - ... But it is not the standard.
 - Do you really need it?
-

1 Minute Example: Serverless Monte Carlo

Simple case (Monte Carlo, not my code) that runs on Google Cloud Functions (FaaS) and returns the value of the calculation.

Up and running in less than 1 minute.

Code:^{*1}

```
# main.py: Numerical Integration using Monte Carlo
method
# FB - 201006137
import math
import random

# define any function here!
def f(x):
    return math.sin(x)

# define any xmin-xmax interval here! (xmin < xmax)
def montecarlo(request):
    xmin = 0.0
    xmax = 2.0 * math.pi

    # find ymin-ymax
    numSteps = 1000000 # bigger the better but slower!
    ymin = f(xmin)
    ymax = ymin
```

```
for i in range(numSteps):
    x = xmin + (xmax - xmin) * float(i) / numSteps
    y = f(x)
    if y < ymin: ymin = y
    if y > ymax: ymax = y

# Monte Carlo
rectArea = (xmax - xmin) * (ymax - ymin)
numPoints = 1000000 # bigger the better but slower!
ctr = 0
for j in range(numPoints):
    x = xmin + (xmax - xmin) * random.random()
    y = ymin + (ymax - ymin) * random.random()
    if math.fabs(y) <= math.fabs(f(x)):
        if f(x) > 0 and y > 0 and y <= f(x):
            ctr += 1 # area over x-axis is positive
        if f(x) < 0 and y < 0 and y >= f(x):
            ctr -= 1 # area under x-axis is negative

fnArea = rectArea * float(ctr) / numPoints
return "Numerical integration = " + str(fnArea)
```

^{*1}: Not my code, numerical recipe from the Internet.

1 Minute Example: Deploying and Running

```
# Deploying (I am logged in google cloud, no other steps needed)  
gcloud functions deploy montecarlo --runtime python37 --trigger-http  
  
# Running (triggering) the code  
curl "https://us-central1-myexperiment.cloudfunctions.net/montecarlo"  
Numerical integration = 0.0024630086404143978
```

HPC for the Cloud - Pitfalls

- Embrace the Order in Chaos: Do NOT fight the nature of the Cloud.
 - Plan, architect and plan again.
 - Be prepared for failure: checking points, incremental backoff, different AZs...
-

Choosing a Provider

- Know your needs, requirements and budget.
 - Follow a methodology (e.g. RACI).
 - Public, Private or Hybrid?
-

Money...

- The Cloud can be expensive... but also cheap... (as well as DCs).
 - Many waivers and programs available for funding (e.g. EU Horizon 2020). Talk to your provider.
-

Takeaways

- The Cloud is not a silver bullet.
 - Understand the paradigm first: read a lot or reach Engineers with expertise.
 - Design and architect are the most important stages.
-

— Thank You!

Thank You! - Questions?



References

- *The NIST Definition of Cloud Computing*. P Mell, et al. NIST, 2011.
 - *Designing Data-Intensive Applications*, Chapter 8. M Kleppmann. O'Reilly, 2017.
 - *Views on the potential of Cloud computing and IaaS/HPaaS for meteorology and climatology*. D Montes, et al. EMS Annual Meeting Abstracts 14 (EMS2017), 639, 2017.
 - *CNCF Cloud Native Interactive Landscape*. Cloud Native Computing Foundation, <<https://landscape.cncf.io>> . October 2017.
 - *Design and Implementation of a Cloud Computing Adoption Decision Tool: Generating a Cloud Road*. I Bidosola, et al. PLoS ONE, 2017.
-